HisabDo — Department 3 Ingestion Execution Report

Prepared 10 November 2026 | Scope: supplied knowledge_base.json and related audit/code evidence

Execution outcome

Staging normalization completed; production ingestion not executed. The supplied knowledge base contains 8 records, all retained in a standardized JSON structure. All 8 remain quarantined because approval, owner, privacy review and/or verification evidence is incomplete. Therefore 0 records were released for production indexing.

What was performed

• Read the supplied knowledge_base.json and retained all 8 source content strings without rewriting their meaning.

• Normalized document IDs, titles, sections, categories, language labels, product-module labels, aliases and audit-state fields into a consistent schema.

• Stored the original content body separately from query aliases; generated SHA-256 hashes for body integrity checks.

• Kept missing source URLs, owners, last-updated dates and module assignments null or explicitly unverified rather than inventing values.

• Prepared a master checklist and type-aware chunking/embedding recommendations.

What was not performed / blockers

• No source crawling or page fetching was run; no official source guide, feature guide, accounting table or workflow document was supplied.

• No production chunking or embedding run was performed, and no vectors were written to an index.

• No exact model commit SHA, dependency lock versions, tokenizer revision, ANN index type or full vector-store parameters were present in supplied evidence.

• No new/modified/deleted-source sync was executed. The supplied code evidence does not establish a complete crawler and delete reconciliation pipeline.

• No feedback/review dataset was supplied; no anonymization or PII-removal run was performed.

• No automated test suite was executed during this document preparation.

Observed technical facts from supplied project materials

• The project context identifies sentence-transformers/all-MiniLM-L6-v2 and a similarity threshold of 0.42 in semantic_search.py; an immutable model revision was not supplied.

• The reviewed semantic-search logic enforces required approval/source metadata and rejects duplicate document IDs. It does not by itself establish duplicate-content detection, page-fetch failure capture, PII anonymization or scheduled-review expiry.

• The historical REPORT.md states 11 tests passed, 0 failed, but notes that the positive retrieval test used an isolated fixture. This historical result was not independently rerun here.

Recommended ingestion run sequence

1. Register source URL, owner, version, source hash, approval and privacy clearance.

2. Fetch source; log HTTP/fetch failures with URL, timestamp, status and retry outcome.

3. Extract content; remove only confirmed navigation/chrome/promotional boilerplate while preserving meaning and retaining a raw-source snapshot.

4. Normalize into the schema; compute content hash; detect duplicate IDs and duplicate content separately.

5. Chunk by content type using the proposed configuration; validate chunk boundaries and metadata inheritance.

6. Load a pinned embedding model revision and record dependency versions, dimension and preprocessing configuration.

7. Upsert new/changed chunks; tombstone/delete chunks for removed source versions; keep an auditable source-to-chunk manifest.

8. Record parse/vector errors, failed records, duplicate decisions and counts; run retrieval regression tests before release.

Release decision

The normalized set is a **staging sample batch**, not a cleaned approved production knowledge set. Keep all 8 records out of production retrieval until genuine source approval, ownership, privacy review, verification dates and authoritative source evidence are provided. Chunking values are engineering recommendations requiring benchmark validation, not measured optimal settings.